

Parallel Algorithm Design

Term 2021 (Winter)

Exercise 1

- Submit electronically as one tar(or tar.gz) file including C++ source files, CMake build system files and an report(PDF) to Moodle until Thursday, 04.11.2021 14:00
- Include names on the top of the sheets.
- A maximum of three students is allowed to work jointly on the exercises.

Reading assignment

G. Amdahl, “[Validity of the Single Processor Approach to Achieving Large-Scale Computing Capabilities](#),” AFIPS Conference Proceedings, 30:483–485, 1967.

[Amdahl’s Law on Wikipedia](#).

J. Gustafson, “[Reevaluating Amdahl’s Law](#),” Communications of the ACM, 31(5), pp. 532–533, 1988.

[Gustafson–Barsis’s Law on Wikipedia](#).

1.1 Allocation and copying (25 points)

Implement the function

```
std::vector<int> matVec(  
    const std::vector<int>& A, const std::vector<int>& x);
```

that multiplies matrix A (row-major storage) with vector x and returns the locally created result y. Because the STL vector supports move semantics, no unnecessary copies are made (be sure to compile with at least C++11 support). Do not assume A to be a square matrix; from the sizes of the `std::vectors` A and x you should be able to derive the number of columns and rows in A.

The function is convenient to use, but it invariably allocates memory for the result, which may not be optimal. Implement a second variant which permits the caller to preallocate a vector:

```
void matVec(  
    std::vector<int>& y,  
    const std::vector<int>& A, const std::vector<int>& x);
```

In these benchmarks, what do you measure primarily? Draw a graph based on the benchmark output obtained with ‘`-benchmark_output csv`’.

1.2 Lambdas (25 points)

Lambdas are a very convenient way of specifying function objects. Let us say we want to still combine **A** and **x** into a new vector **y**, but instead of multiplication we want to use the xor operation, and instead of summation we want to count how many xor-products per row are equal to a certain given mask, e.g. `0x0F`.

Implement a version of `matVec()` that in addition to **A** and **x** takes two function objects to represent the binary operation (addition or xor operation) and the reduction (addition or masked counting). Then in `main()` specify these function objects using lambdas. Test on a small **A**, for example 8×4 , and choose some initial values such that the above mask actually appears for multiple xor-products.

1.3 `std::async()` (40 points)

Now let us further generalize by defining a function `matVecIt()` which does not accept the entire **A**, but rather a begin iterator for the result vector **y** and a begin and an end iterator for the vector **A**, such that we can multiply with a submatrix consisting of some consecutive matrix rows. Having this functionality it becomes easy to implement a parallel version `matVecItPar()`.

`matVecItPar()` takes one more parameter, namely the suggested parallelization factor p . Instead of one call, `matVecItPar()` on a matrix with M rows now performs p asynchronous calls to `matVecIt()` with M/p rows using `std::async()`. Use the matrix from 1.1, check that you get the same result **y**, and compare the execution time for $p = 1, 2, 4, 8, 16$.